

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df=sns.load_dataset("titanic")
df.head()
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked
class \								
0	0	3	male	22.0	1	0	7.2500	S
Third								
1	1	1	female	38.0	1	0	71.2833	C
First								
2	1	3	female	26.0	0	0	7.9250	S
Third								
3	1	1	female	35.0	1	0	53.1000	S
First								
4	0	3	male	35.0	0	0	8.0500	S
Third								

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True

```
df.describe()
```

	survived	pclass	age	sibsp	parch
fare					
count	891.000000	891.000000	714.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594
std	0.486592	0.836071	14.526497	1.102743	0.806057
min	0.000000	1.000000	0.420000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000
50%	0.000000	3.000000	28.000000	0.000000	0.000000
75%	1.000000	3.000000	38.000000	1.000000	0.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000

```
df.shape
```

```
(891, 15)
```

```
df.isnull().sum()
```

```
survived      0
pclass        0
sex           0
age          177
sibsp         0
parch         0
fare          0
embarked       2
class         0
who           0
adult_male    0
deck         688
embark_town    2
alive         0
alone         0
dtype: int64
```

```
df["age"].fillna(df["age"].mean(),inplace=True)
df.isnull().any()
```

```
survived      False
pclass        False
sex           False
age           False
sibsp         False
parch         False
fare          False
embarked      True
class         False
who           False
adult_male    False
deck          True
embark_town   True
alive         False
alone         False
dtype: bool
```

```
df.deck.value_counts()
```

```
deck
C    59
B    47
D    33
E    32
A    15
F    13
```

```

G      4
Name: count, dtype: int64

df["deck"].fillna(df["deck"].mode()[0],inplace=True)
df.deck.value_counts()

deck
C      747
B       47
D       33
E       32
A       15
F       13
G        4
Name: count, dtype: int64

df["embarked"].fillna(df["embarked"].mode()[0],inplace=True)
df.embarked.value_counts()

embarked
S      646
C      168
Q       77
Name: count, dtype: int64

df["embark_town"].fillna(df["embark_town"].mode()[0],inplace=True)
df["embark_town"].value_counts()

embark_town
Southampton    646
Cherbourg      168
Queenstown     77
Name: count, dtype: int64

sns.pairplot(data=df)

<__array_function__ internals>:200: RuntimeWarning: Converting input
from bool to <class 'numpy.uint8'> for compatibility.
<__array_function__ internals>:200: RuntimeWarning: Converting input
from bool to <class 'numpy.uint8'> for compatibility.
C:\Users\HP\AppData\Roaming\Python\Python311\site-packages\seaborn\
axisgrid.py:118: UserWarning: The figure layout has changed to tight
    self._figure.tight_layout(*args, **kwargs)

<seaborn.axisgrid.PairGrid at 0x24e1d0cdc10>

```



```
4      0      3    male  35.0      0      0  8.0500      S
Third
```

```

      who  adult_male  deck  embark_town  alive  alone
0     man           True    C  Southampton    no  False
1  woman           False    C   Cherbourg   yes  False
2  woman           False    C  Southampton   yes  True
3  woman           False    C  Southampton   yes  False
4     man           True    C  Southampton    no  True
```

```
df.drop(['alive','class','who','embark_town'],axis=1)
```

```

      survived  pclass    sex      age  sibsp  parch      fare
embarked \
0           0      3    male  22.000000      1      0   7.2500
S
1           1      1  female  38.000000      1      0  71.2833
C
2           1      3  female  26.000000      0      0   7.9250
S
3           1      1  female  35.000000      1      0  53.1000
S
4           0      3    male  35.000000      0      0   8.0500
S
```

```

..      ...      ...      ...      ...      ...      ...      .
..
886           0      2    male  27.000000      0      0  13.0000
S
887           1      1  female  19.000000      0      0  30.0000
S
888           0      3  female  29.699118      1      2  23.4500
S
889           1      1    male  26.000000      0      0  30.0000
C
890           0      3    male  32.000000      0      0   7.7500
Q
```

```

      adult_male  deck  alone
0           True    C  False
1           False    C  False
2           False    C   True
3           False    C  False
4           True    C   True
..      ...      ...      ...
886          True    C   True
887          False    B   True
888          False    C  False
889          True    C   True
890          True    C   True
```

```
[891 rows x 11 columns]
```


```
corr=df.corr(numeric_only=True)
corr
```

	survived	pclass	age	sibsp	parch	fare
\ survived	1.000000	-0.338481	-0.069809	-0.035322	0.081629	0.257307
pclass	-0.338481	1.000000	-0.331339	0.083081	0.018443	-0.549500
age	-0.069809	-0.331339	1.000000	-0.232625	-0.179191	0.091566
sibsp	-0.035322	0.083081	-0.232625	1.000000	0.414838	0.159651
parch	0.081629	0.018443	-0.179191	0.414838	1.000000	0.216225
fare	0.257307	-0.549500	0.091566	0.159651	0.216225	1.000000
adult_male	-0.557080	0.094035	0.253236	-0.253586	-0.349943	-0.182024
alone	-0.203367	0.135207	0.179775	-0.584471	-0.583398	-0.271832

	adult_male	alone
survived	-0.557080	-0.203367
pclass	0.094035	0.135207
age	0.253236	0.179775
sibsp	-0.253586	-0.584471
parch	-0.349943	-0.583398
fare	-0.182024	-0.271832
adult_male	1.000000	0.404744
alone	0.404744	1.000000


```
sns.heatmap(data=corr,annot=True,)
```

```
<Axes: >
```




```
sns.scatterplot(x="age", y="fare",data=df)
```

```
<Axes: xlabel='age', ylabel='fare'>
```





```
sns.boxplot(data=df, x='pclass', y='age')  
<Axes: xlabel='pclass', ylabel='age'>
```

```
# Box plot for detecting outliers in a variable (e.g., 'Age')
sns.boxplot(x=df['age'])
plt.title('Box Plot of Age')
plt.show()
# Scatter plot for detecting outliers between 'Fare' and 'Age'
sns.scatterplot(data=df, x='age', y='fare')
plt.title('Scatter Plot of Age vs. Fare')
plt.show()
```

Box Plot of Age





```
q1 = df.age.quantile(0.25)
q3 = df.age.quantile(0.75)

IQR=q3-q1
IQR

13.0

upper_limit = q3+1.5*IQR
upper_limit

54.5

lower_limit = q1-1.5*IQR
lower_limit


2.5

df.median(numeric_only=True)
```


survived	0.000000
pclass	3.000000
age	29.699118
sibsp	0.000000
parch	0.000000

```
fare      14.454200
adult_male 1.000000
alone      1.000000
dtype: float64

df['age']=np.where(df['age']>upper_limit,30,df['age'])
sns.boxplot(df.age)
<Axes: >
```



```
sns.boxplot(df.pclass)
<Axes: >
```



```
X=df.drop(columns=["embark_town","alive","alone","class","who","deck"],axis=1)
```

```
X.head()
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked
adult_male								
0	0	3	male	22.0	1	0	7.2500	S
True								
1	1	1	female	38.0	1	0	71.2833	C
False								
2	1	3	female	26.0	0	0	7.9250	S
False								
3	1	1	female	35.0	1	0	53.1000	S
False								
4	0	3	male	35.0	0	0	8.0500	S
True								

```
X=df.drop(columns=["survived"],axis=1)
```

```
y=df["survived"]
```

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
```

```
X["sex"]=le.fit_transform(X["sex"])
X["embarked"]=le.fit_transform(X["embarked"])
X["adult_male"]=le.fit_transform(X["adult_male"])
```

```
X.head()
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked
adult_male								
0	0	3	1	22.0	1	0	7.2500	2
1								
1	1	1	0	38.0	1	0	71.2833	0
0								
2	1	3	0	26.0	0	0	7.9250	2
0								
3	1	1	0	35.0	1	0	53.1000	2
0								
4	0	3	1	35.0	0	0	8.0500	2
1								

```
mapping=dict(zip(le.classes_, range(len(le.classes_))))
mapping
```

```
{False: 0, True: 1}
```

```
from sklearn.preprocessing import MinMaxScaler
ms=MinMaxScaler()
```

```
X_Scaled=ms.fit_transform(X)
X_Scaled=pd.DataFrame(ms.fit_transform(X), columns=X.columns)
X_Scaled.head()
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked
0	0.0	1.0	1.0	0.402762	0.125	0.0	0.014151	1.0
1	1.0	0.0	0.0	0.701381	0.125	0.0	0.139136	0.0
2	1.0	1.0	0.0	0.477417	0.000	0.0	0.015469	1.0
3	1.0	0.0	0.0	0.645390	0.125	0.0	0.103644	1.0
4	0.0	1.0	1.0	0.645390	0.000	0.0	0.015713	1.0

	adult_male
0	1.0
1	0.0
2	0.0
3	0.0
4	1.0

```

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(X_Scaled,y,test_size
=0.2,random_state =0)

print(x_train.shape,x_test.shape,y_train.shape,y_test.shape)

(712, 9) (179, 9) (712,) (179,)

from sklearn.linear_model import LinearRegression
lr=LinearRegression()

lr.fit(x_train,y_train)

LinearRegression()

y_pred=lr.predict(x_test)
y_pred
array([1.08619750e-15, 9.62392253e-16, 5.77631802e-16, 1.00000000e+00,
        1.00000000e+00, 1.00000000e+00, 1.00000000e+00, 1.00000000e+00,
        1.00000000e+00, 1.00000000e+00, 8.77963058e-16, 1.00000000e+00,
        1.02621963e-15, 1.00000000e+00, 1.00000000e+00, 6.63062415e-16,
        9.13481986e-16, 8.61050706e-16, 1.00001432e-15, 1.00000000e+00,
        8.04696537e-16, 1.00000000e+00, 1.02622702e-15, 8.05173570e-16,
        7.19114240e-16, 1.00000000e+00, 9.78646191e-16, 1.00000000e+00,
        1.00000000e+00, 6.35851065e-16, 1.11323716e-15, 1.00000000e+00,
        9.49084893e-16, 1.00000000e+00, 8.99900546e-16, 1.00000000e+00,
        1.05818559e-15, 8.76819140e-16, 8.53881110e-16, 1.06072990e-15,
        1.00000000e+00, 1.05686251e-15, 9.58281490e-16, 5.57510445e-16,
        1.00000000e+00, 9.62187720e-16, 9.62187720e-16, 1.00000000e+00,
        9.54024687e-16, 1.00005480e-15, 1.00000000e+00, 1.00000000e+00,
        1.00000000e+00, 8.85655726e-16, 1.00000000e+00, 6.95365532e-16,
        8.09957887e-16, 6.40443783e-16, 7.15583995e-16, 1.00000000e+00,
        1.03540069e-15, 1.25023893e-15, 1.00000000e+00, 9.27978215e-16,
        1.00000000e+00, 1.04267810e-15, 1.00000000e+00, 9.85435820e-16,
        1.00000000e+00, 1.00000000e+00, 1.00000000e+00, 1.00000000e+00,
        8.32855536e-16, 9.62096514e-16, 9.24906640e-16, 1.00000000e+00,
        6.61901606e-16, 1.00868786e-15, 9.93780483e-16, 1.01417325e-15,
        9.58730350e-16, 1.00000000e+00, 7.30628515e-16, 8.64267839e-16,
        9.96407278e-16, 1.00000000e+00, 1.00000000e+00, 1.00000000e+00,
        1.00000000e+00, 7.45514711e-16, 8.23509348e-16, 8.85655726e-16,
        1.00000000e+00, 1.00000000e+00, 6.85922853e-16, 1.09047334e-15,
        1.00000000e+00, 7.36160996e-16, 8.09957887e-16, 1.00000000e+00,
        6.17380264e-16, 9.95843117e-16, 8.85801111e-16, 9.42663063e-16,
        8.32952200e-16, 1.00000000e+00, 1.00000000e+00, 1.44699233e-15,
        7.88402851e-16, 1.00000000e+00, 1.09312060e-15, 1.00000000e+00,
        1.00000000e+00, 7.75244161e-16, 1.00000000e+00, 1.00000000e+00,
        1.00000000e+00, 1.00000000e+00, 1.03587109e-15, 1.00000000e+00,
        1.00000000e+00, 1.02580069e-15, 8.92355354e-16, 8.82886495e-16,
        9.85350728e-16, 9.35986080e-16, 1.09047583e-15, 1.00000000e+00,
        8.95183625e-16, 1.04014299e-15, 7.77141507e-16, 8.99870746e-16,

```

```

9.29074673e-16, 6.47026346e-16, 8.10283199e-16, 8.99870746e-16,
9.68430568e-16, 7.90295616e-16, 9.24311991e-16, 8.39574885e-16,
8.75340448e-16, 1.00000000e+00, 9.62096514e-16, 7.56273646e-16,
1.00000000e+00, 7.34004549e-16, 8.89646673e-16, 1.00000000e+00,
1.00000000e+00, 9.61734235e-16, 8.25795703e-16, 1.00000000e+00,
8.26558799e-16, 9.57258303e-16, 1.00000000e+00, 8.62512914e-16,
1.13072141e-15, 1.00000000e+00, 8.39367276e-16, 1.00000000e+00,
9.62187720e-16, 1.00000000e+00, 1.00000000e+00, 1.00000000e+00,
8.85120418e-16, 9.28719787e-16, 1.11505312e-15, 9.07174212e-16,
9.21421290e-16, 8.70896016e-16, 1.11446413e-15, 1.00000000e+00,
9.62569696e-16, 9.58281490e-16, 1.00000000e+00, 9.62096514e-16,
1.00000000e+00, 9.14292782e-16, 8.85564521e-16])

```

y_test

```

495    0
648    0
278    0
31     1
255    1
..
780    1
837    0
215    1
833    0
372    0

```

Name: survived, Length: 179, dtype: int64

```

survived=pd.DataFrame({"actual_survived":y_test,"Predicted
_survived":y_pred})
survived

```

	actual_survived	Predicted _survived
495	0	1.086197e-15
648	0	9.623923e-16
278	0	5.776318e-16
31	1	1.000000e+00
255	1	1.000000e+00
..
780	1	1.000000e+00
837	0	9.620965e-16
215	1	1.000000e+00
833	0	9.142928e-16
372	0	8.855645e-16

[179 rows x 2 columns]

```

from sklearn import metrics

```

```

print(metrics.r2_score(y_test,y_pred))

```


1.0

```
print(metrics.mean_squared_error(y_test,y_pred))
```

1.4836608224492802e-30

```
print(np.sqrt(metrics.mean_squared_error(y_test,y_pred)))
```

1.218056165556121e-15